



Lower Bounds

CS240

Spring 2020

Rui Fan



Upper and lower bounds

- What is the minimum resources (time, space, etc.) needed to solve a problem?
- Consider sorting n numbers.
 - Insertion sort takes $O(n^2)$ time.
 - This puts an upper bound of $O(n^2)$ on the time to sort n numbers.
 - Merge sort takes $O(n \log n)$ time.
 - This puts an upper bound of $O(n \log n)$ on the time to sort n numbers.
- We want to make the upper bound as low as possible, i.e. solve the problem faster.
- Suppose an algorithm A solves problem X in $f(n)$ time when input size is x .
 - Then $f(n)$ is an upper bound on the complexity of X .



Upper and lower bounds

- What about the least amount of time to solve X ?
- Suppose we know that any algorithm that solves X takes at least $g(n)$ time, when X has size n .
 - Then $g(n)$ is a lower bound on the complexity of X .
- If the lower bound $g(n)$ is large, it means problem X is hard to solve.
 - **Ex** NP-Hard problems are hard because they (probably) have super-exponential lower bounds.
- To show a lower bound, we need to give a proof.
 - Usually we show if an algorithm takes too little time, it must sometimes produce the wrong answer.
- The lower bound for a problem depends on the computational model.
 - If a model has very powerful primitive operations, then algorithms can run faster, and the lower bound is smaller.
- If the complexity of an algorithm for problem X matches the lower bound for problem X , the algorithm is optimal, and the lower bound is tight.



A warm-up

- Say we want to find the larger of two numbers x and y .
 - We can do this with 1 comparison, so this is an upper bound.
 - What's the lower bound? Do we need at least 1 comparison? Can we do 0 comparisons?
 - No. Suppose an algorithm doesn't compare x and y .
 - So basically, the algorithm declares either x or y to be bigger, without looking at them.
 - Say the algorithm declares x bigger. Then let's set $y > x$.
 - Algorithm won't notice this, cause it doesn't compare x and y .
 - So algorithm still declares x is bigger, which is wrong.
 - This type of argument is called indistinguishability, and is frequently used when proving lower bounds.
 - Same argument if algorithm always declares y bigger without comparing.
 - Hence, any algorithm must do at least 1 comparison, so 1 is a lower bound.



Outline

- We'll prove lower bounds for the following problems.
 - Merging two lists.
 - Finding the max.
 - Finding the max and min.
 - Sorting n numbers.



Merging two lists

- How many comparisons needed to merge two lists of size n into sorted order?
 - During the execution, the algorithm can compare some input elements a and b , and get back response “ $a < b$ ”, “ $a = b$ ” or “ $a > b$ ”.
- If the lists are sorted, $2n-1$ comparisons is an upper bound.
- Let's prove this is also a lower bound.
- Let the input lists be a_1, a_2, \dots, a_n and b_1, b_2, \dots, b_n , and suppose $a_1 < b_1 < a_2 < b_2 < \dots < a_n < b_n$.
 - So the algorithm must output $a_1, b_1, a_2, b_2, \dots, a_n, b_n$.
- When comparing some a_i and b_j , it gets back the following response:
 - $a_i < b_j$ if $i \leq j$.
 - $a_i > b_j$ if $i > j$.
- We show any algorithm has to perform $\geq 2n - 1$ comparisons to merge the two lists.
 - This gives a $2n-1$ time lower bound on merging, since a merging algorithm must correctly merge any two input lists, including the two lists above.

Merging two lists

- **Claim** Any correct algorithm must compare a_i to b_i , for every i .
 - Suppose not; say the algorithm doesn't compare a_1 to b_1 .
 - Now, if the input was actually $b_1 < a_1 < a_2 < b_2 < \dots < a_n < b_n$, then the algorithm still outputs $a_1, b_1, \dots, a_n, b_n$, which is wrong.
 - Because the algorithm doesn't compare a_1 and b_1 , it can't distinguish the new input from the original.
 - Same argument if algorithm doesn't compare a_i to b_i , for any i .
 - So algorithm does n comparisons of this type.
- **Claim** Any correct algorithm must compare b_i to a_{i+1} , for every $i < n$.
 - If not, then say it doesn't compare b_1 to a_2 . Then it can't distinguish original input from input $a_1 < a_2 < b_1 < b_2 < \dots < a_n < b_n$, and will give wrong answer.
 - Thus, $n-1$ comparisons of this type.
- So, any algorithm must do at least $2n-1$ comparisons.
- So $2n-1$ is a lower bound on the complexity to merge into sorted order.



Finding the max

- How many comparisons to find the largest number in an unsorted array of n distinct numbers.
- Upper bound: $n-1$.
- Lower bound: also $n-1$.
- To prove this, we'll keep track of what information the algorithm learns as it executes.
 - Say algorithm never compared some element to any other element.
 - Then the algorithm doesn't know anything about this element. It could be the max, or not the max.
 - Thus, the algorithm can't correctly output the max without comparing this element to some others.
 - Say there are two elements, and both are larger than every element they've been compared to.
 - Then either one of them could be the max.
 - So algorithm can't output the max without comparing these two elts.
- Let's formalize this intuition.



Finding the max

- At any stage of the alg, give every array element one of 3 colors, white, blue or red.
 - White means this element has never been compared to any other element.
 - Blue means this element is bigger than all the elements it's been compared to.
 - Red means this element was smaller than some element it was compared to.
 - Let w_k, b_k, r_k be number of white, blue and red elements after A has done k comparisons.
 - So initially, $w_0=n$ and $b_0=r_0=0$.
- We'll show that for any k , $w_k+b_k \geq n-k$.
- We'll show that as long as $w_k+b_k > 1$, A can't terminate.
- Hence, when A terminates, we have $w_k+b_k = 1$, and A must have done $k \geq n-1$ comparisons.

Finding the max

- **Claim** For any k , $w_k + b_k \geq n - k$.
- **Proof** By induction on k . Claim holds for $k=0$.
 - For larger k , consider the k 'th comparison. It must either be between: 2 white elements (WW case), a white and blue element (WB case), a white and red (WR), 2 reds (RR), 2 blues (BB), red and blue (RB).
 - Do a case by case analysis.
 - **WW**: Make the first element $>$ second element.
 - This is possible, because both elements are white, so neither have been in any comparisons, so they can be in either order.
 - After comparison, first element becomes blue, second element red.
 - Number of whites decreases by 2, blues increases by 1.
 - **By induction**, $w_{k-1} + b_{k-1} \geq n - k + 1$. Also, $w_k = w_{k-1} - 2$, and $b_k = b_{k-1} + 1$. So $w_k + b_k \geq n - k$.
 - **WB**: Make the first element $<$ second element.
 - This is possible, since first element hasn't been in any comparisons.
 - So first element becomes red, second remains blue.
 - So $w_k = w_{k-1} - 1$, $b_k = b_{k-1}$, so $w_k + b_k \geq n - k$.



Finding the max

- **WR**: Make the first element $>$ second element. First element becomes blue, second stays red.
 - So $w_k = w_{k-1} - 1$, $b_k = b_{k-1} + 1$, so $w_k + b_k \geq n - k + 1 > n - k$.
- **RR**: Make first element $>$ second element. Both elements stay red.
 - $w_k + b_k = w_{k-1} + b_{k-1} \geq n - k + 1 > n - k$.
- **BB**: Make first element $>$ second element. First one stays blue, second becomes red.
 - $w_k + b_k = w_{k-1} + b_{k-1} - 1 \geq n - k$.
- **RB**: Make first element $<$ second element. Both elements stay same color.
 - $w_k + b_k = w_{k-1} + b_{k-1} \geq n - k + 1 > n - k$.
- Hence $w_k + b_k \geq n - k$ by induction.



Finding the max

- **Claim** Suppose after making k comparisons, we have $w_k + b_k > 1$. Then A cannot terminate.
- **Proof** Say A terminates, and outputs a value x as the max.
 - Since $w_k + b_k > 1$, either $w_k \geq 1$, or $b_k > 1$.
 - If $w_k \geq 1$, then there's a white element y that's never been compared to x (or any other elt).
 - Make $y > x$. Then the algorithm is wrong.
 - If $b_k > 1$, then there are at least 2 blue elements.
 - x must be a blue element.
 - If x is red, it's not max.
 - x is not white, by above.
 - Take another blue element z . x and z were never compared.
 - If they had been, either x or z would have turned red.
 - Make $z > x$. Now A is wrong.
- Since A can't terminate as long as $w_k + b_k > 1$, then $k \geq n-1$ when A terminates.
- So A does $\geq n-1$ comparisons.



Finding the max and min

- How many comparisons does it take to find the max and min elements in an **unsorted array A of n distinct numbers**.
- Upper bound 1: $2n-2$ comparisons.
- Upper bound 2: $3n/2-2$ comparisons.
 - Pair up the elements, $A[1]$ and $A[2]$, $A[3]$ and $A[4]$, etc.
 - Compare the elements in each pair ($n/2$ comps total).
 - Put all the bigger elements in a temp array Big, put all the smaller elements in temp array Small.
 - Big and Small each have size $n/2$.
 - Find the max element in Big and output it as max of A ($n/2-1$ comparisons).
 - Find the min element in Small and output it as the min of A ($n/2-1$ comparisons).
- Upper bound 3: $3n/2-2$ comparisons via divide and conquer.
 - Exercise.
- Lower bound: $3n/2-2$ comparisons!



Finding the max and min

- Intuition for proof is similar to one for max.
- At any stage of alg, give each array element one of 4 colors, white, blue, red and purple, representing what the algorithm knows about the element.
 - White means this element has never been compared against any other element.
 - Blue means this element is bigger than all the elements it's been compared against.
 - Red means this element was smaller than every element it was compared against.
 - Purple means this element was bigger than some elt(s) it was compared to, and smaller than some other(s).



Finding the max and min

- To terminate, algorithm must eliminate all white elements, since these could be the min or max.
- Also, algorithm can only leave one blue and one red.
 - Else either of two blues can be max, either of two reds can be min.
- As comparisons happen, algorithm gets more info, and elements change color, e.g. from white to blue, red to purple, etc.
- Too few comparisons means the algorithm doesn't have time to eliminate all whites, and all but 1 blue and red.
- Proof keeps track of number of whites, blues and reds after some number of comparisons.



Finding the max and min

- Label each comparison by its type.
 - E.g. WW is comparison between two white elts.
 - There are 10 types, WW, WB, WR, WP, BB, BR, BP, RR, RP, PP.
- Denote the number of comparisons of type WW by ww , number of WB comps by wb , etc. 10 numbers total.
- Let w , b , r denote number of whites, blues and reds, resp., at some stage of the algorithm.



Finding the max and min

- **Claim 1** When A terminates, $w=0$ and $b=r=1$.
- **Proof** Say A outputs x as max, y as min.
 - Neither x nor y can be white, since we can make a white element be neither max nor min.
 - If there is a white element z when A terminates, we can make $z > x$, and A is wrong. So $w=0$.
 - x must be a blue element, as in the finding max proof.
 - If there's another blue element z , then x and z weren't compared, so we can make $z > x$, and A is wrong. So $b=1$.
 - y must be a red element.
 - If there's another red element z , then we can make $z < y$, and A is wrong. So $r=1$.

Finding the max and min

- The table states what happens when each type of comparison occurs. Similar to the case analysis in finding max proof.
 - **Ex** If WW occurs, make the first element $>$ second element (denoted $E_1 > E_2$), so these elements become blue and red (BR).
 - **Ex** If WB occurs, we make the first element $<$ second element (denoted $E_1 < E_2$), so the elements become red and blue (RB).

Comparison type	Result	Comparison type	Result
WW	$E_1 > E_2$, BR	BR	$E_1 > E_2$, BR
WB	$E_1 < E_2$, RB	BP	$E_1 > E_2$, BP
WR	$E_1 > E_2$, BR	RR	$E_1 < E_2$, RP
WP	$E_1 > E_2$, BP	RP	$E_1 < E_2$, RP
BB	$E_1 > E_2$, BP	PP	$E_1 < E_2$, PP



Finding the max and min

- **Claim 2** At any stage of the alg, we have

- $w = n - 2ww - rw - bw - pw.$

- $b = ww + rw + pw - bb.$

- $r = ww + bw - rr.$

- **Proof** These follow just by counting w,b,r using the table on the previous page.

- For w, there are initially n whites. Each WW comparison removes 2 whites. Each RW, BW or PW comp removes 1 white.

- For b, each WW, RW or PW comparison creates 1 blue element. Each BB comparison removes 1 blue.

- For r, each WW, BW comparison creates 1 red element. Each RR removes 1 red.

Finding the max and min

- **Theorem** Any algorithm performs at least $3n/2 - 2$ comparisons.
- **Proof** The total number of comparisons is $C = ww + wb + wr + wp + bb + br + bp + rr + rp + pp$.
 - By claims 1 and 2, when A terminates we have $2ww + rw + bw + pw = n$, $bb = ww + rw + pw - 1$, $rr = ww + bw - 1$.
 - So $bb + rr = 2ww + rw + bw + pw - 2 = n - 2$.
 - $C \geq ww + wb + wr + wp + bb + rr$
 $= ww + wb + wr + wp + n - 2$
 $= n - ww + n - 2$
 $= 2n - 2 - ww$.
 - $ww \leq n/2$, because each WW comp decreases number of whites by 2, and there are only n whites.
 - So $C \geq 3n/2 - 2$.



Sorting

- How many comparisons are needed to sort n numbers?
- Upper bound: $O(n \log n)$ using merge sort.
- Lower bound: $\Omega(n \log n)$.
- To prove the lower bound, we first need a model for how a comparison-based sorting algorithm works.
 - This is called the decision tree model.
- The lower bound is not valid in other models.
 - If an algorithm can do things besides comparing two numbers, e.g. look at the digits of a number, it can sort faster than $\Omega(n \log n)$ time.
 - Lower bounds can be very sensitive to the computational model.




Decision trees

- In this model, in each step, algorithm can only compare a pair of numbers x , y .
- Based on result of the comparison, it decides next pair of numbers to compare.
 - So an execution of the algorithm is a sequence of comparisons, each comparison determined by result of previous comparison.
- When the algorithm terminates, it outputs a permutation representing the sorted order of the input.
- The complexity of the algorithm is the most number of comparisons it does before terminating.

Decision trees

- Model behavior of the algorithm by a binary tree.
 - Each internal node is a pair of number x,y to compare.
 - If $x \leq y$, go to left child. If $x > y$, go to right child.
 - Each leaf represents an output, and is labeled with a permutation representing the sorted order of the inputs.
- An execution is simply a path from root to a leaf.
 - At any node, the algorithm has obtained some info from the comparisons it's done.
 - It uses this info to decide the next comparison to do.
 - Eventually, it obtains enough info to generate an output.
- Complexity of algorithm is the length of the longest root-leaf path.



Lower bound for sorting

- Given n numbers as input, they can be in $n!$ different orders.
- Given an input order, algorithm must output that order.
 - So decision tree of algorithm must have a leaf labeled with that order.
 - So the decision tree **has $\geq n!$ leafs.**
- Say height of decision tree is h .
 - The complexity of the algorithm is h .
 - Since decision tree is binary, it has $\leq 2^h$ leaves.
- So $2^h \geq (\# \text{ leaves of dec tree}) \geq n!$, and so $h \geq \log_2(n!)$.
 - $\log_2(n!) = \log_2 n + \log_2(n-1) + \dots + \log_2 1 \geq \log_2 n + \log_2(n-1) + \dots + \log_2(n/2) \geq \frac{n}{2}(\log_2 n - 1) = \Omega(n \log n)$.
 - Can also use Stirling's approximation.
- So we proved the algorithm does $\Omega(n \log n)$ comparisons.